

Grounding LLMs in a Formal Ontology

A pervasive knowledge-graph binding that makes every model measurably smarter

F1 0.37 → 0.81 across 5 LLMs

VisionFlow ■ VisionClaw ■ Agentbox

<http://www.visionflow.info>



The headline

Connecting a **formal ontology** (4,196 OWL classes, 222k inferred axioms) to *every* AI call lifts factual recall ***across the board***.

+0.44 mean F1

Augmented **0.81** vs. ungrounded **0.37** ■ hallucination roughly **halved**

5 models ■ 16 KG-grounded questions ■ 160 isolated runs ■ objective scoring

Lead, not buried: the binding works, and it is model-agnostic.



What we built


- **VisionFlow** — the immersive 3D knowledge-graph + agent platform (visionflow.info).
- **VisionClaw** — the Rust engine: Oxigraph/Whelk ontology store, GPU physics, real-time graph.
- **Agentbox** — the sovereign agent runtime; 100+ skills, MCP tooling, governed memory.

The final piece: a **pervasive ontology binding** so any AI call can ground itself in the formal knowledge graph — **read-pervasive**, write-governed, budget-bounded, **fail-open**.

Features are legion; this deck leads with the one that ties them together.



The binding, in one picture



- **Read-pervasive:** every agent, consultant and turn can consult the KG.
- **Write-governed:** proposals are auth-gated and queued; derived facts are fenced.


How we measured it (objectively)



- **KG-as-oracle:** ground truth generated *from the graph itself* — neighbours, subclasses, class existence — so scoring is deterministic, not subjective.
- **Clean A/B:** each cell is an *isolated* session given only the question; augmented arm receives the ontology subgraph, control uses parametric knowledge only.
- **5 models × 16 questions × 2 conditions = 160 isolated runs.**
- **Grader:** precision / recall / F1 + hallucination, token-set matched.


Anthropic Opus/Sonnet/Haiku, Google Gemini 2.5 Pro, Z.AI GLM-5.2.

Result: every model wins



Universal lift: +0.31 to +0.54 F1. The smallest model (Haiku) gains the most.


Result: hallucination roughly halved



Grounding replaces plausible-but-wrong guesses with the graph's actual vocabulary.




Where grounding helps most



Biggest gains on proprietary structure (subclasses: +0.50)
and niche concepts the base model can't know.

What it costs




- Grounding adds context tokens and one retrieval round-trip — **optional and per-call**.
- Gemini 2.5 Pro: 4198 prompt tokens/query; GLM-5.2: 3620 — modest for the accuracy gained.
- **Budget-bounded & fail-open:** if the graph is unreachable, the turn proceeds ungrounded — never blocked.

Net: a bounded, switchable cost for a large, universal accuracy gain.

The binding augments thinking without overpowering the context window.

The design works



A formal ontology, bound pervasively to AI, makes **every** model more accurate and less hallucinatory.

- +0.44 mean F1 across 5 LLMs; hallucination roughly halved.
- Read-pervasive, write-governed, budget-bounded, fail-open — production-shaped.
- One shared brain across tool, consultant and CLI surfaces.

VisionFlow ▪ VisionClaw ▪ Agentbox

<http://www.visionflow.info>